

TITLE: Fraud Detection System For Online Transaction

Subtitle: Applying RandomForest and XGBoost for Fraud Detection and Developed with Streamlit



Group 3:

**Khansa Rubab
Shannon Chen
Victoria Giles
Zalak Shah**

Table of Contents

1. Introduction
2. Data Acquisition and Exploration
3. Data Preprocessing
4. Model Selection and Evaluation
 - Hyperparameter Tuning
 - Model Performance Metrics
5. Model Development with Streamlit
6. Key Takeaways

INTRODUCTION



Background on the problem of fraud detection:

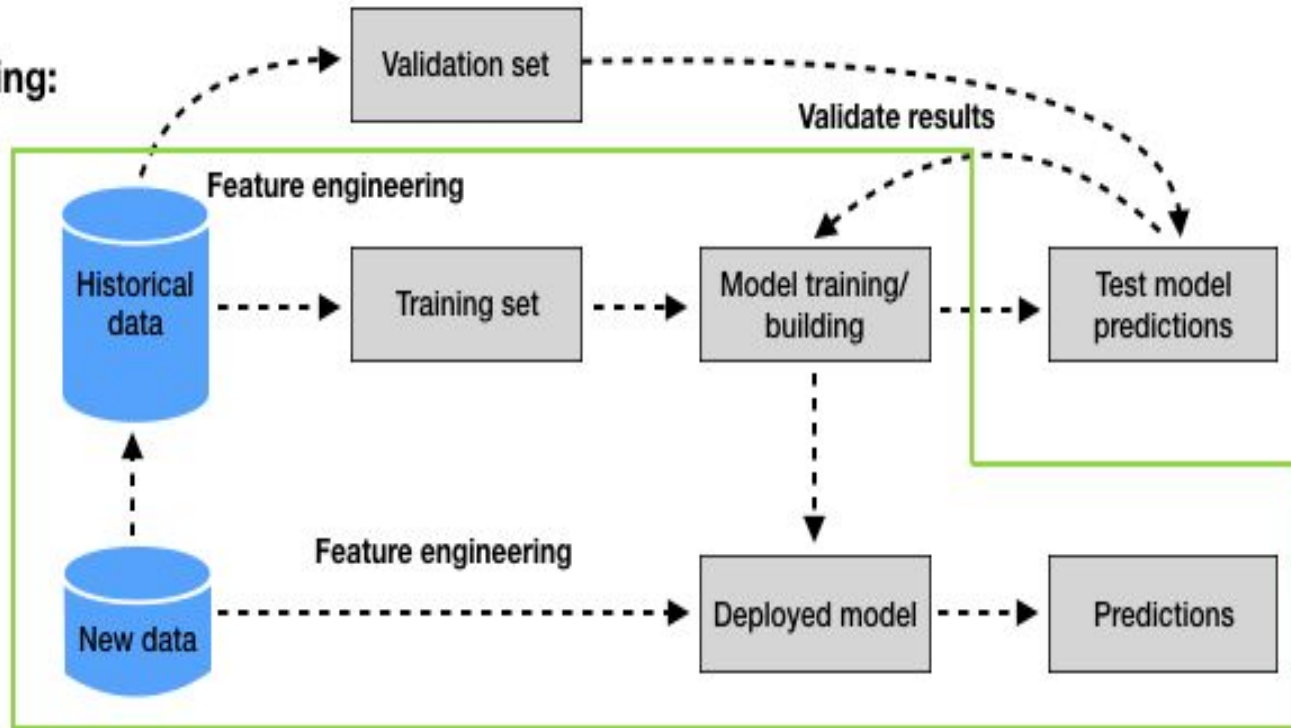
- ❑ Tackling the challenge of online transaction fraud detection.
- ❑ The significance of timely and accurate fraud identification.
- ❑ Goal: Create a machine learning model to assess the probability of fraudulent transactions (classification problem).

Targeted Supervised Learning Solutions for Data-Driven Decision Making:

Supervised learning:

- Feature engineering
- Model training
- Model validation

Production



Dataset:



- Source: Kaggle (<https://www.kaggle.com/competitions/ieee-fraud-detection/data>)
- Consists of transactional and identity data
- Features:** TransactionID, TransactionDT, TransactionAmt, ProductCD, card1-6, addr1-2, dist1-2, P_emaildomain, R_emaildomain, C1-C14, D1-D15, M1-M9, V1-V339, id_01-id_38, DeviceType, DeviceInfo
- Target:** is Fraud



Overcoming Dataset Complexity:



❑ Challenges faced:

- ❑ Dataset size and complexity make hyperparameter tuning difficult on Jupyter and Google Colab.
- ❑ Computational limitations impacted model training and evaluation
- ❑ Aws free tier does not work.


❑ Concurrent attempt to transition to Apache Spark from Pandas

- ❑ Leverage Spark's distributed computing capabilities
- ❑ Efficiently handle large-scale data processing tasks
- ❑ Improve model training and evaluation times
- ❑ **Limitations:** Having to learn a new language on a fly

Data Processing Workflow:

Data Pre-processing:

- ❑ Conduct exploratory data analysis
- ❑ Combine identity and transaction datasets
- ❑ Manage missing values:
 - ❑ Remove columns with over 40% missing values
 - ❑ Impute remaining missing values using respective column means
 - ❑ Replace infinity values with NaN and fill with column mean
- ❑ Save cleaned data to CSV file
- ❑ Encode categorical variables with LabelEncoder
- ❑ Standardize numerical features using StandardScaler
- ❑ Store processed data in an SQL database



Data exploration and cleaning process

Step 2: Exploring the Data

```
# get shape of the data
print("test_identity Shape: ", test_identity.shape)
print("test_transaction Shape: ", test_transaction.shape)
print("train_identity Shape: ", train_identity.shape)
print("train_transaction Shape: ", train_transaction.shape)
```

```
test_identity Shape: (141907, 41)
test_transaction Shape: (506691, 393)
train_identity Shape: (144233, 41)
train_transaction Shape: (590540, 394)
```

```
# print first two rows of each dataset
print(test_identity.head())
print(test_transaction.head())
print(train_identity.head())
print(train_transaction.head(2))
```

	TransactionID	id-01	id-02	id-03	id-04	id-05	id-06	id-07	id-08
0	3663586	-45.0	280290.0	NaN	NaN	0.0	0.0	NaN	NaN
1	3663588	0.0	3579.0	0.0	0.0	0.0	0.0	NaN	NaN
2	3663597	-5.0	185210.0	NaN	NaN	1.0	0.0	NaN	NaN
3	3663601	-45.0	252944.0	0.0	0.0	0.0	0.0	NaN	NaN
4	3663602	-95.0	328680.0	NaN	NaN	7.0	-33.0	NaN	NaN

	id-09	...	id-31	id-32	id-33	id-34	\
0	NaN	...	chrome 67.0 for android	NaN	NaN	NaN	
1	0.0	...	chrome 67.0 for android	24.0	1280x720	match_status:2	
2	NaN	...	ie 11.0 for tablet	NaN	NaN	NaN	
3	0.0	...	chrome 67.0 for android	NaN	NaN	NaN	
4	NaN	...	chrome 67.0 for android	NaN	NaN	NaN	

	id-35	id-36	id-37	id-38	DeviceType	DeviceInfo
0	F	F	T	F	mobile	MYA-L13 Build/HUAWEIMYA-L13
1	T	F	T	T	mobile	LGLS676 Build/MXB48T
2	F	T	T	F	desktop	Trident/7.0
3	F	F	T	F	mobile	MYA-L13 Build/HUAWEIMYA-L13
4	F	F	T	F	mobile	SM-G9650 Build/R16NW

In [6]:

```
# get information about the data
print(test_identity.info())
print(test_transaction.info())
print(train_identity.info())
print(train_transaction.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 141907 entries, 0 to 141906
Data columns (total 41 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   TransactionID        141907 non-null  int64
1   id-01               141907 non-null  float64
2   id-02              136976 non-null  float64
3   id-03              66481 non-null   float64
4   id-04              66481 non-null   float64
5   id-05              134750 non-null  float64
6   id-06              134750 non-null  float64
7   id-07              5059 non-null    float64
8   id-08              5059 non-null    float64
9   id-09              74338 non-null   float64
10  id-10              74338 non-null   float64
11  id-11              136778 non-null  float64
12  id-12              141907 non-null  object
13  id-13              130286 non-null  float64
14  id-14              71357 non-null   float64
15  id-15              136977 non-null  object
16  id-16              125747 non-null  object
17  id-17              135966 non-null  float64
18  id-18              50875 non-null   float64
19  id-19              135906 non-null  float64
20  id-20              135633 non-null  float64
21  id-21              5059 non-null    float64
22  id-22              5062 non-null    float64
23  id-23              5062 non-null    object
24  id-24              4740 non-null    float64
25  id-25              5039 non-null    float64
26  id-26              5047 non-null    float64
27  id-27              5062 non-null    object
28  id-28              136778 non-null  object
29  id-29              136778 non-null  object
30  id-30              70659 non-null   object
31  id-31              136625 non-null  object
32  id-32              70671 non-null   float64
33  id-33              70671 non-null   object
34  id-34              72175 non-null   object
```


Data exploration and cleaning process

The train dataset has a target column called `isFraud`.

In [7]:

```
# get descriptive statistics for each dataset
print(test_identity.describe())
print(test_transaction.describe())
print(train_identity.describe())
print(train_transaction.describe())
```

	TransactionID	id-01	id-02	id-03
count	1.419070e+05	141907.000000	136976.000000	66481.000000
mean	3.972166e+06	-11.325734	192658.729909	0.053008
std	1.469966e+05	14.508520	182613.277215	0.684551
min	3.663586e+06	-100.000000	2.000000	-12.000000
25%	3.859268e+06	-12.500000	63339.500000	0.000000
50%	4.001774e+06	-5.000000	133189.500000	0.000000
75%	4.105284e+06	-5.000000	265717.500000	0.000000
max	4.170239e+06	0.000000	999869.000000	11.000000

	id-04	id-05	id-06	id-07	id-08
count	66481.000000	134750.000000	134750.000000	5059.000000	5059.000000
mean	-0.087454	1.246033	-6.803829	12.493180	-36.577782
std	0.840351	5.071394	15.921457	11.678206	25.544185
min	-19.000000	-81.000000	-100.000000	-41.000000	-100.000000
25%	0.000000	0.000000	-6.000000	3.000000	-46.000000
50%	0.000000	0.000000	0.000000	12.000000	-33.000000
75%	0.000000	1.000000	0.000000	21.000000	-23.000000
max	0.000000	52.000000	0.000000	59.000000	0.000000

	id-09	...	id-17	id-18	id-19
count	74338.000000	...	135966.000000	50875.000000	135906.000000
mean	0.076219	...	191.070341	14.795735	350.122982
std	1.009687	...	30.749535	2.318496	139.140824
min	-32.000000	...	100.000000	11.000000	100.000000
25%	0.000000	...	166.000000	13.000000	266.000000
50%	0.000000	...	166.000000	15.000000	321.000000
75%	0.000000	...	225.000000	15.000000	427.000000
max	16.000000	...	228.000000	29.000000	670.000000

	id-20	id-21	id-22	id-24	id-25
count	135633.000000	5059.000000	5062.000000	4740.000000	5039.000000
mean	408.886230	507.727021	15.336823	13.166667	332.043064
std	158.971756	227.371061	5.618032	3.222440	86.356683
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000	0.000000
max	0.000000	0.000000	0.000000	0.000000	0.000000

train_data Shape: (500540, 434)

test_data Shape: (506691, 433)

In [11]:

```
# Check for missing values in train and test data
missing_train = train.isnull().sum().sort_values(ascending=False)
missing_test = test.isnull().sum().sort_values(ascending=False)
```

In [12]:

```
# Display the percentage of missing values in each column
print("Missing values in train (%):")
print((missing_train / len(train)) * 100)
print("\nMissing values in test_data (%):")
print((missing_test / len(test)) * 100)
```

Missing values in train (%):

id_24	99.196159
id_25	99.130965
id_07	99.127070
id_08	99.127070
id_21	99.126393
...	...

C11	0.000000
C12	0.000000
C13	0.000000
C14	0.000000
TransactionID	0.000000

Length: 434, dtype: float64

Missing values in test_data (%):

id-24	99.064519
id-25	99.005508
id-26	99.003929
id-07	99.001561
id-08	99.001561
...	...

V111	0.000000
V112	0.000000
V113	0.000000
V114	0.000000
TransactionID	0.000000

Length: 433, dtype: float64

Note:

**drop columns with a missing value percentage greater than a certain threshold (let's say 50%)

Step 5: Handle Missing Values

Data exploration and cleaning process

```
In [13]: # Drop columns with more than 40% missing values
train_data = train.drop(columns=missing_train[missing_train > 0.40 * len(train)].index)
test_data = test.drop(columns=missing_test[missing_test > 0.40 * len(test)].index)
```

```
In [14]: # Impute missing values in the remaining columns with their respective means
train_data.fillna(train_data.mean(), inplace=True)
test_data.fillna(test_data.mean(), inplace=True)
```

```
In [15]: # Replace infinity values with NaN
train_data.replace([np.inf, -np.inf], np.nan, inplace=True)
test_data.replace([np.inf, -np.inf], np.nan, inplace=True)

# Fill NaN values with the mean of each column
train_data.fillna(train_data.mean(), inplace=True)
test_data.fillna(test_data.mean(), inplace=True)
```

```
In [16]: # Load cleaned data
clean_train_data = pd.read_csv("Resources/clean_train_data.csv")
clean_test_data = pd.read_csv("Resources/clean_test_data.csv")
```

```
In [17]: # from sklearn.preprocessing import LabelEncoder


# # Identify categorical columns
# categorical_columns = X_train.select_dtypes(include=['object']).columns

# # Apply Label encoding
# for col in categorical_columns:
#     le = LabelEncoder()
#     le.fit(pd.concat([X_train[col], X_val[col], test_data[col]]))
#     X_train[col] = le.transform(X_train[col])
#     X_val[col] = le.transform(X_val[col])
#     test_data[col] = le.transform(test_data[col])
```

```
In [18]: # X_train.to_csv("Resources/X_train.csv", index=False)
# X_val.to_csv("Resources/X_val.csv", index=False)
# y_train.to_csv("Resources/y_train.csv", index=False)
# y_val.to_csv("Resources/y_val.csv", index=False)
# test_data.to_csv("Resources/test_data.csv", index=False)
```


Data Visualisation

- Examine target variable distribution (isFraud) in the train_transaction dataset



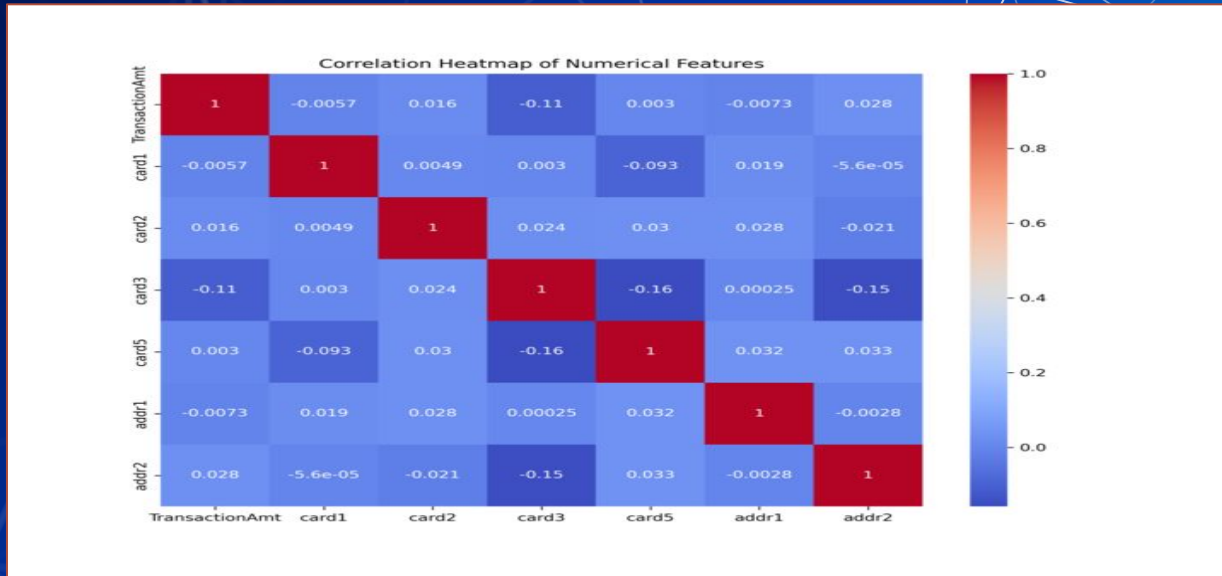
Distribution of the Variable for both:

- Compare Transaction Amt distributions for fraud and non-fraud transactions



Heatmap to visualise correlations

- Investigate correlations between selected numerical features using a heatmap:
- Features: Transaction_Amt, card1, card2, card3, card5, addr1, addr2



- Gain insights into data patterns and relationships to inform model development and feature selection.

Model Selection:

Explore various machine learning models

- ❑ RandomForest
- ❑ Ensemble methods:
 - Extremely Random tree classifier
 - AdaBoost
 - RandomForest
- ❑ Evaluate model performance with classification report

Random Forest Classifier:

Training Score: 0.999778733588467

Testing Score: 0.9801469841162326

	precision	recall	f1-score	support
0	0.98	1.00	0.99	142497
1	0.93	0.47	0.62	5138
accuracy			0.98	147635
macro avg	0.95	0.73	0.80	147635
weighted avg	0.98	0.98	0.98	147635

Extremely Random Trees Classifier:

Training Score: 1.0

Testing Score: 0.9809462525823822

	precision	recall	f1-score	support
0	0.98	1.00	0.99	142497
1	0.92	0.50	0.65	5138
accuracy			0.98	147635
macro avg	0.95	0.75	0.82	147635
weighted avg	0.98	0.98	0.98	147635

AdaBoost Classifier:

Training Score: 0.9705106061119202

Testing Score: 0.9705760829071697

	precision	recall	f1-score	support
0	0.97	1.00	0.98	142497
1	0.82	0.20	0.32	5138
accuracy			0.97	147635
macro avg	0.90	0.60	0.65	147635
weighted avg	0.97	0.97	0.96	147635

Our Findings for Model Selection and Evaluation

- ❑ Identified Random Forest as the best-performing model based on evaluation metrics
- ❑ We initially found Random Forest to be the top performer, but after examining more resources, XGBoost also emerged as a strong contender.
- ❑ Selected Random Forest and XGBoost for in-depth analysis
- ❑ Retrieved data from the SQL database for model training and testing
- ❑ Compared Random Forest and XGBoost performance on the retrieved dataset
- ❑ Finalized the model selection for app development based on the comparison results

Model Evaluation:

- ❑ Assess models performance using ROC AUC score and classification report
- ❑ Compare RandomForest and XGBoost models
- ❑ Identify the best-performing model based on evaluation metrics

Decision: Developed the XGBoost model for real-time fraud detection with streamlit

Random Forest:

	precision	recall	f1-score	support
0	0.98	1.00	0.99	113866
1	0.94	0.47	0.63	4242
accuracy			0.98	118108
macro avg	0.96	0.74	0.81	118108
weighted avg	0.98	0.98	0.98	118108

Accuracy: 0.9800352220002032

ROC AUC: 0.7362497890665184

XGBoost:

	precision	recall	f1-score	support
0	0.98	1.00	0.99	113866
1	0.92	0.48	0.63	4242
accuracy			0.98	118108
macro avg	0.95	0.74	0.81	118108
weighted avg	0.98	0.98	0.98	118108

Accuracy: 0.9798574186337928

ROC AUC: 0.7394484317086844

Model Optimisation (Attempted)

Attempted with RandomForest & XGBoost Classifiers

Utilised hyperparameter tuning technique: *GridSearchCV*

```
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier
```

```
# Define the parameter grids for RandomForest and XGBoost
rf_param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5],
    'min_samples_leaf': [1, 2]
}
```

```
from sklearn.model_selection import GridSearchCV
from xgboost import XGBClassifier
```

```
# Define the parameter grids for XGBoost

xgb_param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [6, 10, 15],
    'learning_rate': [0.01, 0.1],
    'subsample': [0.5, 1],
    'colsample_bytree': [0.5, 1]
}
```

It performs an exhaustive search over all possible combinations of hyperparameters, training and evaluating the model with each combination using cross validation.

- ❑ Returns the hyperparameters that resulted in best performance.

Model Optimisation (Timed out)

Insufficient computing power despite trying various mediums:

- ❑ Pandas on Jupyter Notebook
- ❑ Pandas on Google Colab (with TPU)
- ❑ PySpark on Google Colab (with TPU)

Thus, unable to successfully implement GridSearchCV as a form of hyperparameter tuning.

However, the **XGBoost Classification model** that we tested has some built-in regularization techniques to improve model generalisation, such as:

- ❑ L1, L2
- ❑ Max depth constraints



Your PC ran into a problem and needs to restart. We're just collecting some error info, and then we'll restart for you. (0% complete)

Model Deployment (Streamlit)

Decision taken to deploy our best performing model (XGBoost Classifier) using 2 features (out of our initial 400 + in the raw dataset) on Streamlit.io.

Created a web app with input fields for each feature in a respective dataset.

Users can upload a new pre-processed dataset onto the app and use it to predict the probability that a transaction is fraudulent.

```
import streamlit as st
import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler
from xgboost import XGBClassifier

# Load the preprocessed data
X_train = pd.read_csv("Resources/X_train.csv")
y_train = pd.read_csv("Resources/y_train.csv")["isFraud"]

# Select two basic features
selected_features = ['TransactionDT', 'TransactionAmt']
X_train_selected = X_train[selected_features]

# Train the XGBoost model
xgb_model = XGBClassifier(use_label_encoder=False, random_state=42)
xgb_model.fit(X_train_selected, y_train, eval_metric='logloss')

# Streamlit app
st.title("Fraud Detection")

transaction_dt = st.number_input("TransactionDT", min_value=0, value=100000)
transaction_amt = st.number_input("TransactionAmt", min_value=0.0, value=50.0)

if st.button("Predict"):
    input_data = pd.DataFrame({"TransactionDT": [transaction_dt],
                              "TransactionAmt": [transaction_amt]})
    prediction = xgb_model.predict(input_data)
    st.write("Prediction: ", "Fraud" if prediction[0] == 1 else "Not Fraud")
```

Model Deployment (Streamlit)

Limitation of current app:

It only takes in two basic features for predictions. It was vital that we had a functioning front-end given the tight timeframe.

Future refinements:

Enhance model performance by adding more features, incorporating other ML models, utilising alternative hyperparameter tuning techniques, and deploying the improved model on Heroku for increased accessibility.

Fraud Detection

TransactionDT

100000

TransactionAmt

50.00

Predict

Prediction: Not Fraud

Key TakeAways:

- ❑ Request more instances of fraudulent transactions from the organization to improve model training and prediction accuracy.
- ❑ Address imbalanced dataset challenges with techniques such as oversampling, undersampling, or using cost-sensitive learning.
- ❑ Optimize computational resources and training times by leveraging scalable solutions like Apache Spark or distributed computing.
- ❑ Seek clarity on column names to potentially merge related variables, inform feature engineering, and improve model interpretability.



Github Link : https://github.com/rubab-malik/Project_4



Thank you Mortaza and Jeffery